

CMPINF 2100 Week 09

Review PCA with the Sonar data set

Import Modules

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

import seaborn as sns
```

Also import functions from sklearn.

```
In [2]: from sklearn.preprocessing import StandardScaler
```

```
In [3]: from sklearn.decomposition import PCA
```

Read data

```
In [4]: sonar_url = 'https://archive.ics.uci.edu/ml/machine-learning-databases/undocumented/connectionist-bench/sonar/sonar.all'
```

```
In [5]: sonar_df = pd.read_csv( sonar_url, header=None )
```

```
In [6]: sonar_df.shape
```

```
Out[6]: (208, 61)
```

```
In [7]: sonar_df.columns
```

```
Out[7]: Int64Index([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15, 16,
                  17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33,
                  34, 35, 36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49, 50,
                  51, 52, 53, 54, 55, 56, 57, 58, 59, 60],
                  dtype='int64')
```

Let's rename the COLUMNS to show the following pattern:

X01, X02, X03 so and and so on.

```
In [8]: sonar_df.columns = ['X%02d' % d for d in sonar_df.columns ]
```

```
In [9]: sonar_df.columns
```

```
Out[9]: Index(['X00', 'X01', 'X02', 'X03', 'X04', 'X05', 'X06', 'X07', 'X08', 'X09',  
          'X10', 'X11', 'X12', 'X13', 'X14', 'X15', 'X16', 'X17', 'X18', 'X19',  
          'X20', 'X21', 'X22', 'X23', 'X24', 'X25', 'X26', 'X27', 'X28', 'X29',  
          'X30', 'X31', 'X32', 'X33', 'X34', 'X35', 'X36', 'X37', 'X38', 'X39',  
          'X40', 'X41', 'X42', 'X43', 'X44', 'X45', 'X46', 'X47', 'X48', 'X49',  
          'X50', 'X51', 'X52', 'X53', 'X54', 'X55', 'X56', 'X57', 'X58', 'X59',  
          'X60'],  
          dtype='object')
```

```
In [10]: sonar_df.dtypes.value_counts()
```

```
Out[10]: float64    60  
object         1  
dtype: int64
```

```
In [11]: sonar_df.select_dtypes('object').info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 208 entries, 0 to 207  
Data columns (total 1 columns):  
#   Column  Non-Null Count  Dtype  
---  -  
0   X60      208 non-null     object  
dtypes: object(1)  
memory usage: 1.8+ KB
```

```
In [12]: sonar_df.X60.value_counts()
```


```
Out[12]: M    111  
         R     97  
         Name: X60, dtype: int64
```

The numeric columns are HIGHLY correlated!!!

```
In [13]: fig, ax = plt.subplots()
```

```
sns.heatmap( data = sonar_df.corr(numeric_only=True),
             vmin=-1, vmax=1, center=0,
             cmap='coolwarm',
             ax=ax)
```

```
plt.show()
```



PCA exploits correlation!!!

PCA crafts or creates new variables based on correlated numeric columns!!!

Extract or Select the numeric columns

```
In [14]: sonar_features = sonar_df.select_dtypes('number').copy()
```

```
In [17]: sonar_features.shape
```

```
Out[17]: (208, 60)
```

Standardize the numeric features

```
In [15]: Xsonar = StandardScaler().fit_transform( sonar_features )
```

```
In [16]: Xsonar.shape
```

```
Out[16]: (208, 60)
```

PCA

BUT...we will NOT use the same approach that we have used up to this point.

We will **NOT** specify the `n_components` argument to `PCA()`.

Let's see what happens if we use the DEFAULT arguments to `PCA()` meaning will NOT specify arguments when the PCA object is initialized!!!!

```
In [18]: sonar_pca = PCA().fit_transform( Xsonar )
```

Check the shape...

```
In [19]: sonar_pca.shape
```

```
Out[19]: (208, 60)
```

There are as many columns returned as the number of columns in the data!!!!

We have previously specified `n_components=2` just to support simple visualization.

But...PCA is capable of giving you MORE than just 2 new variables.

PCA creates as many NEW variables as there are in the original data!!!!

```
In [20]: type( sonar_pca )
```

```
Out[20]: numpy.ndarray
```

Convert `sonar_pca` to a DataFrame using the naming pattern:

pc01, pc02, pc03, etc...

```
In [21]: ['pc%02d' % d for d in range(1, sonar_pca.shape[1]+1)]
```

```
Out[21]: ['pc01',  
          'pc02',  
          'pc03',  
          'pc04',  
          'pc05',  
          'pc06',  
          'pc07',  
          'pc08',  
          'pc09',  
          'pc10',  
          'pc11',  
          'pc12',  
          'pc13',  
          'pc14',  
          'pc15',  
          'pc16',  
          'pc17',  
          'pc18',  
          'pc19',  
          'pc20',  
          'pc21',  
          'pc22',  
          'pc23',  
          'pc24',  
          'pc25',  
          'pc26',  
          'pc27',  
          'pc28',  
          'pc29',  
          'pc30',  
          'pc31',  
          'pc32',  
          'pc33',  
          'pc34',  
          'pc35',  
          'pc36',  
          'pc37',  
          'pc38',  
          'pc39',  
          'pc40',  
          'pc41',  
          'pc42',  
          'pc43',  
          'pc44',  
          'pc45',
```

```
'pc46',  
'pc47',  
'pc48',  
'pc49',  
'pc50',  
'pc51',  
'pc52',  
'pc53',  
'pc54',  
'pc55',  
'pc56',  
'pc57',  
'pc58',  
'pc59',  
'pc60']
```

Create the DataFrame containing ALL PCs!!!

```
In [22]: sonar_pca_df = pd.DataFrame( sonar_pca,  
                                     columns=['pc%d' % d for d in range(1, sonar_pca.shape[1]+1)])
```

```
In [23]: sonar_pca_df.shape
```

```
Out[23]: (208, 60)
```

```
In [24]: sonar_pca_df.columns
```


```
Out[24]: Index(['pc01', 'pc02', 'pc03', 'pc04', 'pc05', 'pc06', 'pc07', 'pc08', 'pc09',  
               'pc10', 'pc11', 'pc12', 'pc13', 'pc14', 'pc15', 'pc16', 'pc17', 'pc18',  
               'pc19', 'pc20', 'pc21', 'pc22', 'pc23', 'pc24', 'pc25', 'pc26', 'pc27',  
               'pc28', 'pc29', 'pc30', 'pc31', 'pc32', 'pc33', 'pc34', 'pc35', 'pc36',  
               'pc37', 'pc38', 'pc39', 'pc40', 'pc41', 'pc42', 'pc43', 'pc44', 'pc45',  
               'pc46', 'pc47', 'pc48', 'pc49', 'pc50', 'pc51', 'pc52', 'pc53', 'pc54',  
               'pc55', 'pc56', 'pc57', 'pc58', 'pc59', 'pc60'],  
              dtype='object')
```

Visualize the Principal Components

Use Seaborn WIDE FORMAT plotting to examine the BOXPLOT or summary stats for each of the PCs!!!

```
In [26]: sns.catplot(data = sonar_pca_df, kind='box', aspect=3.5)
```

```
plt.show()
```



Let's use the `.describe()` method to look at the values of the summary statistics.

```
In [28]: sonar_pca_df.describe().round(3)
```

```
Out[28]:
```

	pc01	pc02	pc03	pc04	pc05	pc06	pc07	pc08	pc09	pc10	...	pc51	pc52	pc53	pc54	pc55
count	208.000	208.000	208.000	208.000	208.000	208.000	208.000	208.000	208.000	208.000	...	208.000	208.000	208.000	208.000	208.000
mean	0.000	-0.000	-0.000	-0.000	0.000	-0.000	0.000	-0.000	-0.000	0.000	...	0.000	-0.000	0.000	0.000	-0.000
std	3.502	3.375	2.270	1.850	1.737	1.565	1.406	1.355	1.244	1.226	...	0.171	0.169	0.152	0.149	0.139
min	-6.221	-7.168	-5.974	-5.976	-4.127	-4.364	-3.312	-5.386	-3.103	-4.013	...	-0.409	-0.481	-0.404	-0.414	-0.383
25%	-2.394	-2.754	-1.564	-1.119	-1.273	-0.993	-0.890	-0.859	-0.737	-0.690	...	-0.118	-0.103	-0.110	-0.084	-0.107
50%	-0.806	-0.047	-0.255	0.001	-0.146	0.064	-0.068	-0.030	0.034	0.031	...	0.003	0.000	0.009	0.002	0.003
75%	1.456	2.238	1.626	0.808	1.206	1.005	0.931	0.926	0.627	0.694	...	0.119	0.104	0.105	0.089	0.079
max	11.727	8.774	6.224	5.984	7.101	5.103	4.023	6.281	4.673	4.129	...	0.421	0.533	0.413	0.479	0.394

8 rows × 60 columns

Visualize the standard deviation!


```
In [29]: sonar_pca_df.describe().loc['std']
```

```
Out[29]: pc01    3.502415
          pc02    3.375367
          pc03    2.270413
          pc04    1.850399
          pc05    1.737458
          pc06    1.565497
          pc07    1.406023
          pc08    1.355253
          pc09    1.243790
          pc10    1.225513
          pc11    1.118563
          pc12    1.070844
          pc13    1.026280
          pc14    0.963094
          pc15    0.927802
          pc16    0.905826
          pc17    0.862753
          pc18    0.839390
          pc19    0.788323
          pc20    0.768270
          pc21    0.754443
          pc22    0.731511
          pc23    0.711589
          pc24    0.681716
          pc25    0.659735
          pc26    0.647951
          pc27    0.609103
          pc28    0.566128
          pc29    0.562450
          pc30    0.546511
          pc31    0.535819
          pc32    0.512801
          pc33    0.473867
          pc34    0.450554
          pc35    0.432673
          pc36    0.427526
          pc37    0.416476
          pc38    0.383997
          pc39    0.363776
          pc40    0.355919
          pc41    0.333572
          pc42    0.308869
          pc43    0.287138
          pc44    0.275326
          pc45    0.248871
```


```
pc46    0.239218
pc47    0.237829
pc48    0.206556
pc49    0.189185
pc50    0.177284
pc51    0.171398
pc52    0.169086
pc53    0.152103
pc54    0.148690
pc55    0.139306
pc56    0.127642
pc57    0.121896
pc58    0.107941
pc59    0.106226
pc60    0.081477
Name: std, dtype: float64
```

Use Pandas plotting methods to show the standard deviation vs the PC number.

```
In [31]: fig, ax = plt.subplots()

sonar_pca_df.describe().loc[ 'std' ].plot(ax=ax)
ax.set_xlabel('Principal Component')
ax.set_ylabel('Standard deviation')

plt.show()
```



The LOW NUMBERED PCs have GREATER variation than the HIGHER NUMBERED PCs!!!!

This is by design!

PCA CREATES the new variables such that PC01 has the HIGHEST variation.

Then, PCA creates PC02 to have the NEXT highest variation.

Then, PCA creates PC03 to have the NEXT highest variation.

So on and so on, the variation decreases for each subsequent PC!!!!

The maximum number of PCs equals the number of columns in the data!!!!

One more important aspect of the PCs!!!


We previously saw that the original 60 numeric columns were correlated.

Let's check the correlation structure of the newly created PCs.

```
In [33]: fig, ax = plt.subplots()

sns.heatmap(data = sonar_pca_df.corr(numeric_only=True),
            vmin=-1, vmax=1, center=0,
            cmap='coolwarm',
            ax=ax)

plt.show()
```



PCA is created such that the NEW variables are UNCORRELATED!!!!!!


PCA is created such that the LOW NUMBERED PCs have the HIGHEST variation while the HIGHEST numbered PCs have the LOWEST VARIATION!!!

But ALL PCs are UNCORRELATED!!!!

Why does this matter?

We have used the FIRST TWO PCs to help our visualizations.

```
In [34]: sns.pairplot(data = sonar_pca_df.iloc[:, :2])  
plt.show()
```

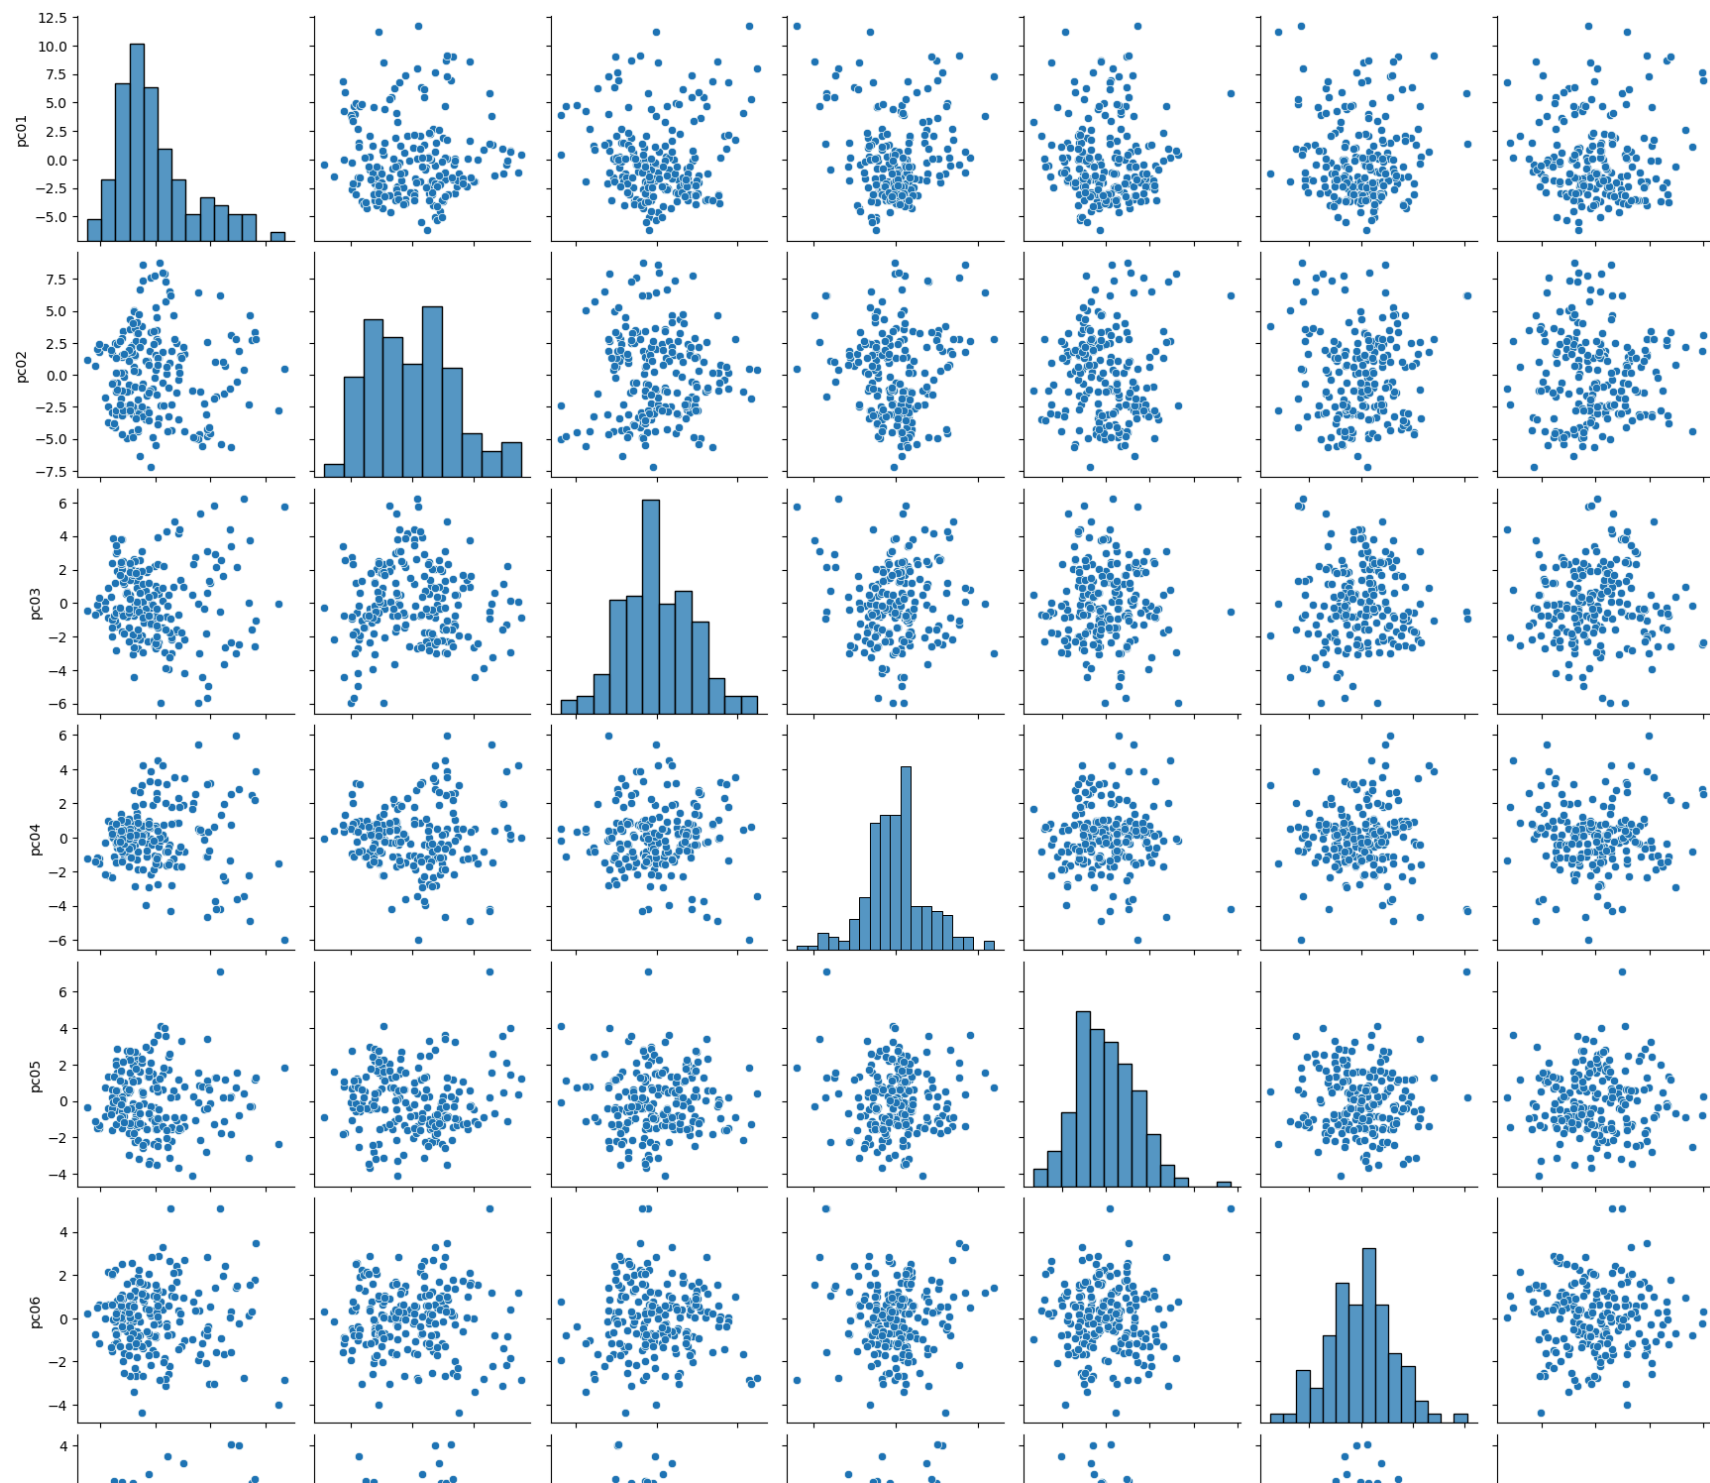



But...we now know there are MORE PCs!!!!

We could consider exploring MORE PCs than just the first 2!!!

This is helpful when there are dozens if not hundreds of variables in the data and those variables have some kind of correlation structure!

```
In [35]: sns.pairplot(data = sonar_pca_df.iloc[:, :7])  
plt.show()
```





You can also group or condition PCs by categorical variables!

```
In [36]: sonar_pca_df['X60'] = sonar_df.X60
```

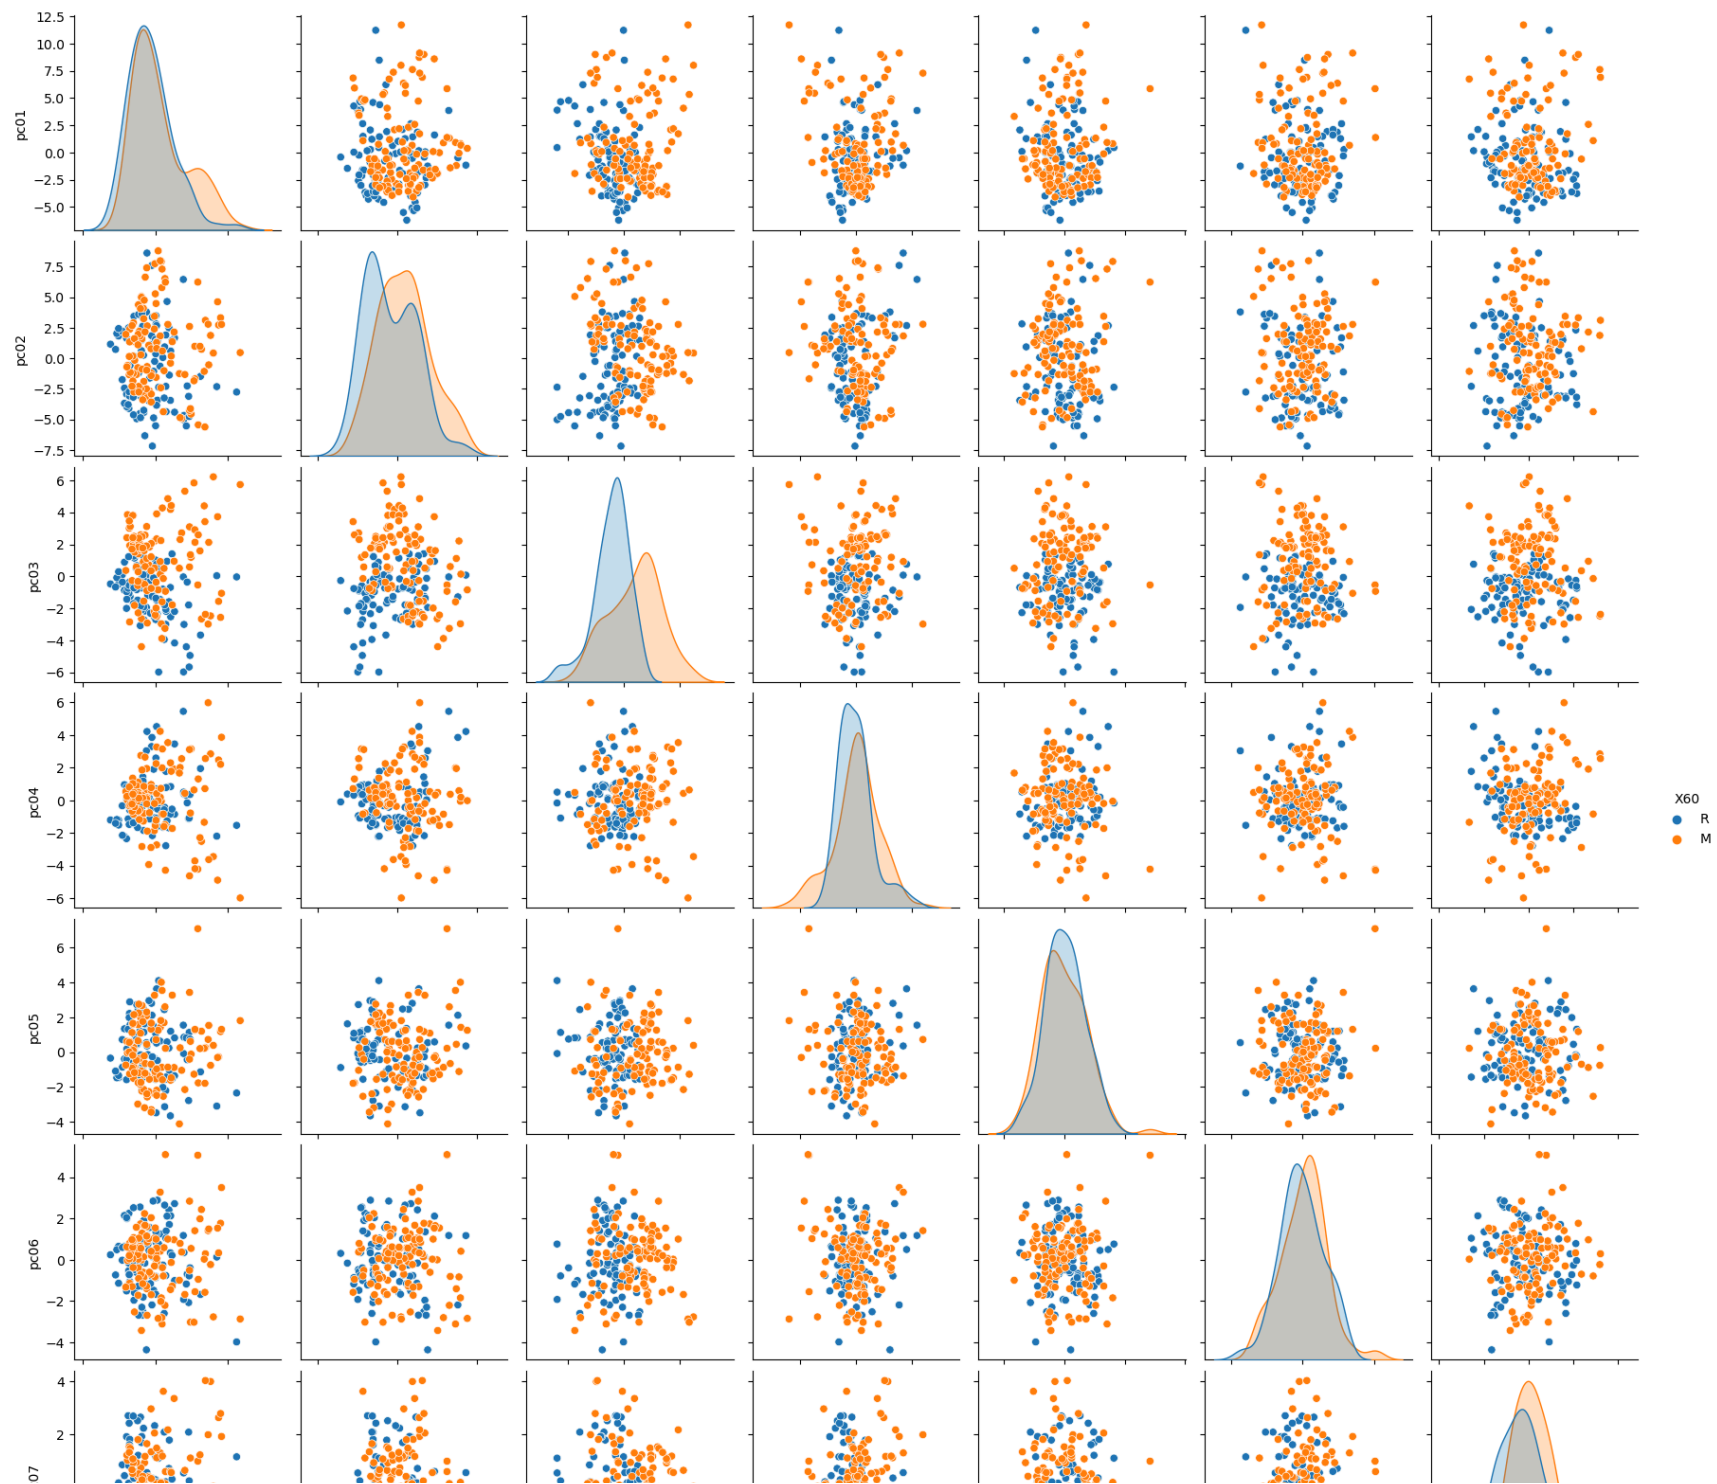
```
In [37]: sonar_pca_df.dtypes
```


```
Out[37]: pc01    float64
pc02    float64
pc03    float64
pc04    float64
pc05    float64
...
pc57    float64
pc58    float64
pc59    float64
pc60    float64
X60      object
Length: 61, dtype: object
```

Let's first use pairs plot to show the CONDITIONAL KDE and conditional scatter plot between the PAIRS of PCs GIVEN X60.

```
In [38]: sns.pairplot(data = sonar_pca_df,
                    vars=['pc%02d' % d for d in range(1, 7+1)],
                    hue='X60',
                    diag_kws={'common_norm': False})

plt.show()
```





In []: